

Introduction to Dynamic Memory Management

FACULTY OF
ENGINEERING

COMP2017/COMP9017



THE UNIVERSITY OF
SYDNEY



Memory

- Memory is a long array of 8 bit pieces called *bytes*
- This array is indexed from 0 to the number of bytes in the memory
- Each index is a memory *address*

0 1 2 3



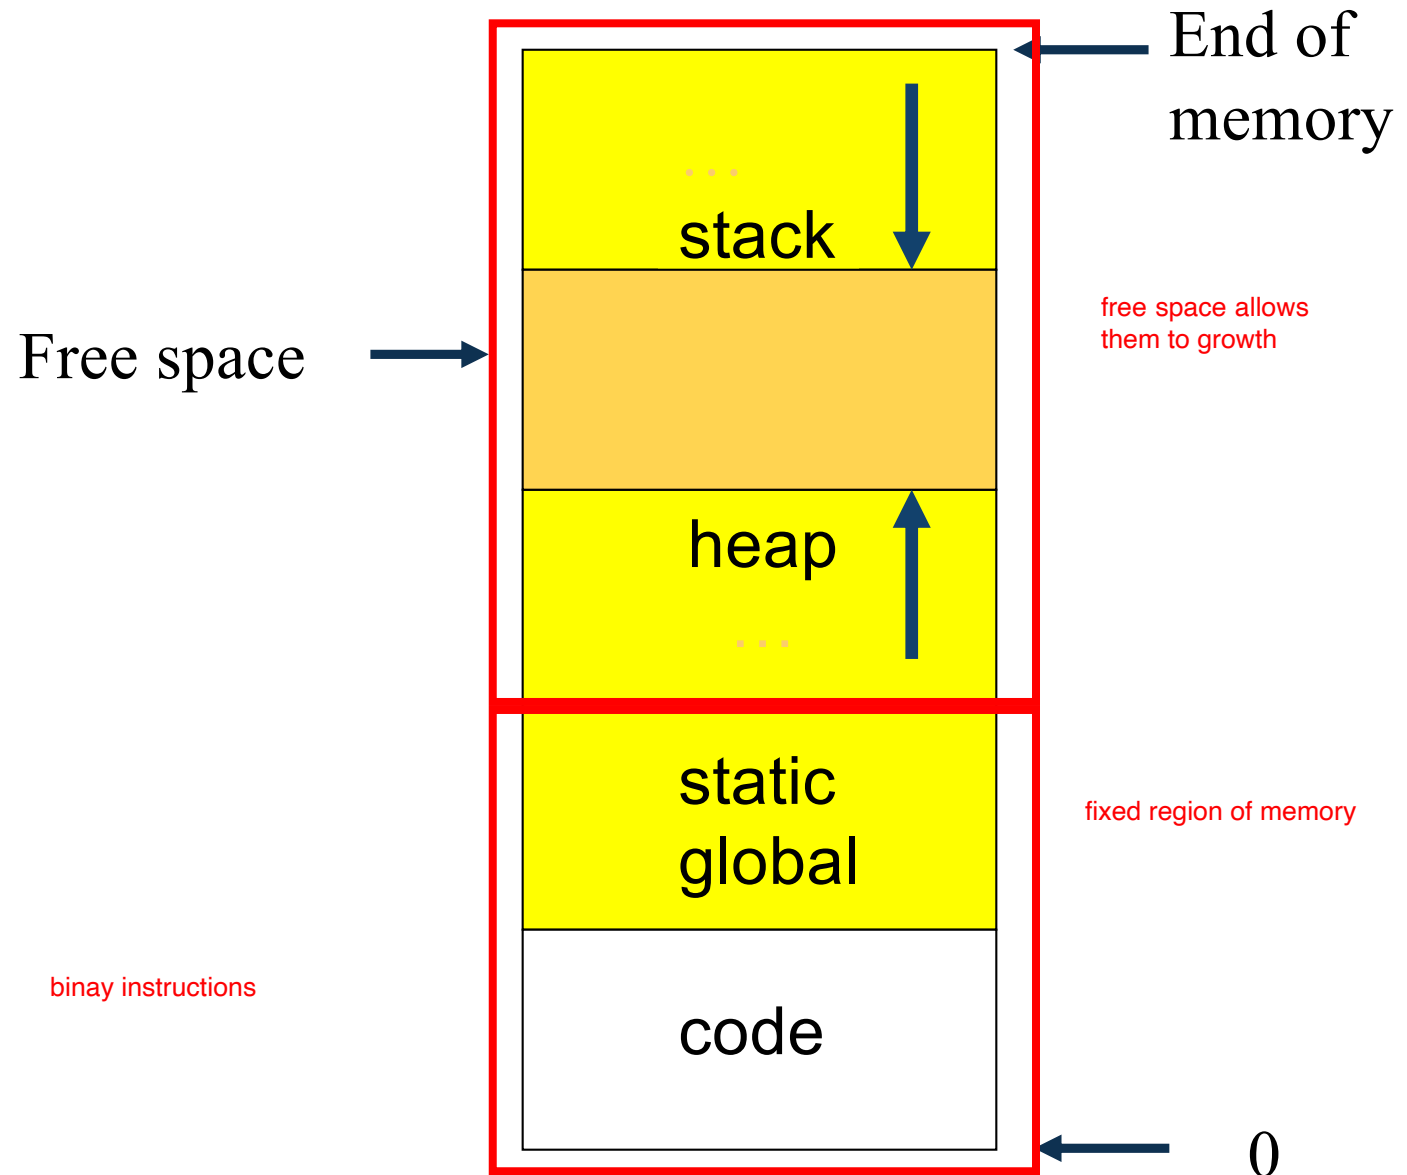
Memory Areas

stack is used for functions call, and the area of
stack memory is used at run time

- Stack: local variables, function arguments, return addresses, temporary storage
- Heap: dynamically allocated memory
- Global/static: global variables, static variables
- Code: program instructions




Memory Layout





The Stack



main
parameter of function
address when return back
local variables in functions

- In C, all variables local to a function and function arguments are stored on the stack
- To call a function the code does:

push arguments onto stack

push return address onto stack

jump to function code

return program counter (next instruction
to execute)



The Stack

- Inside the function, the code does the following:

increment the stack pointer to allow
space for the local variables

execute the code


pop local variables and arguments off the stack

push the return result onto the stack

jump to return address

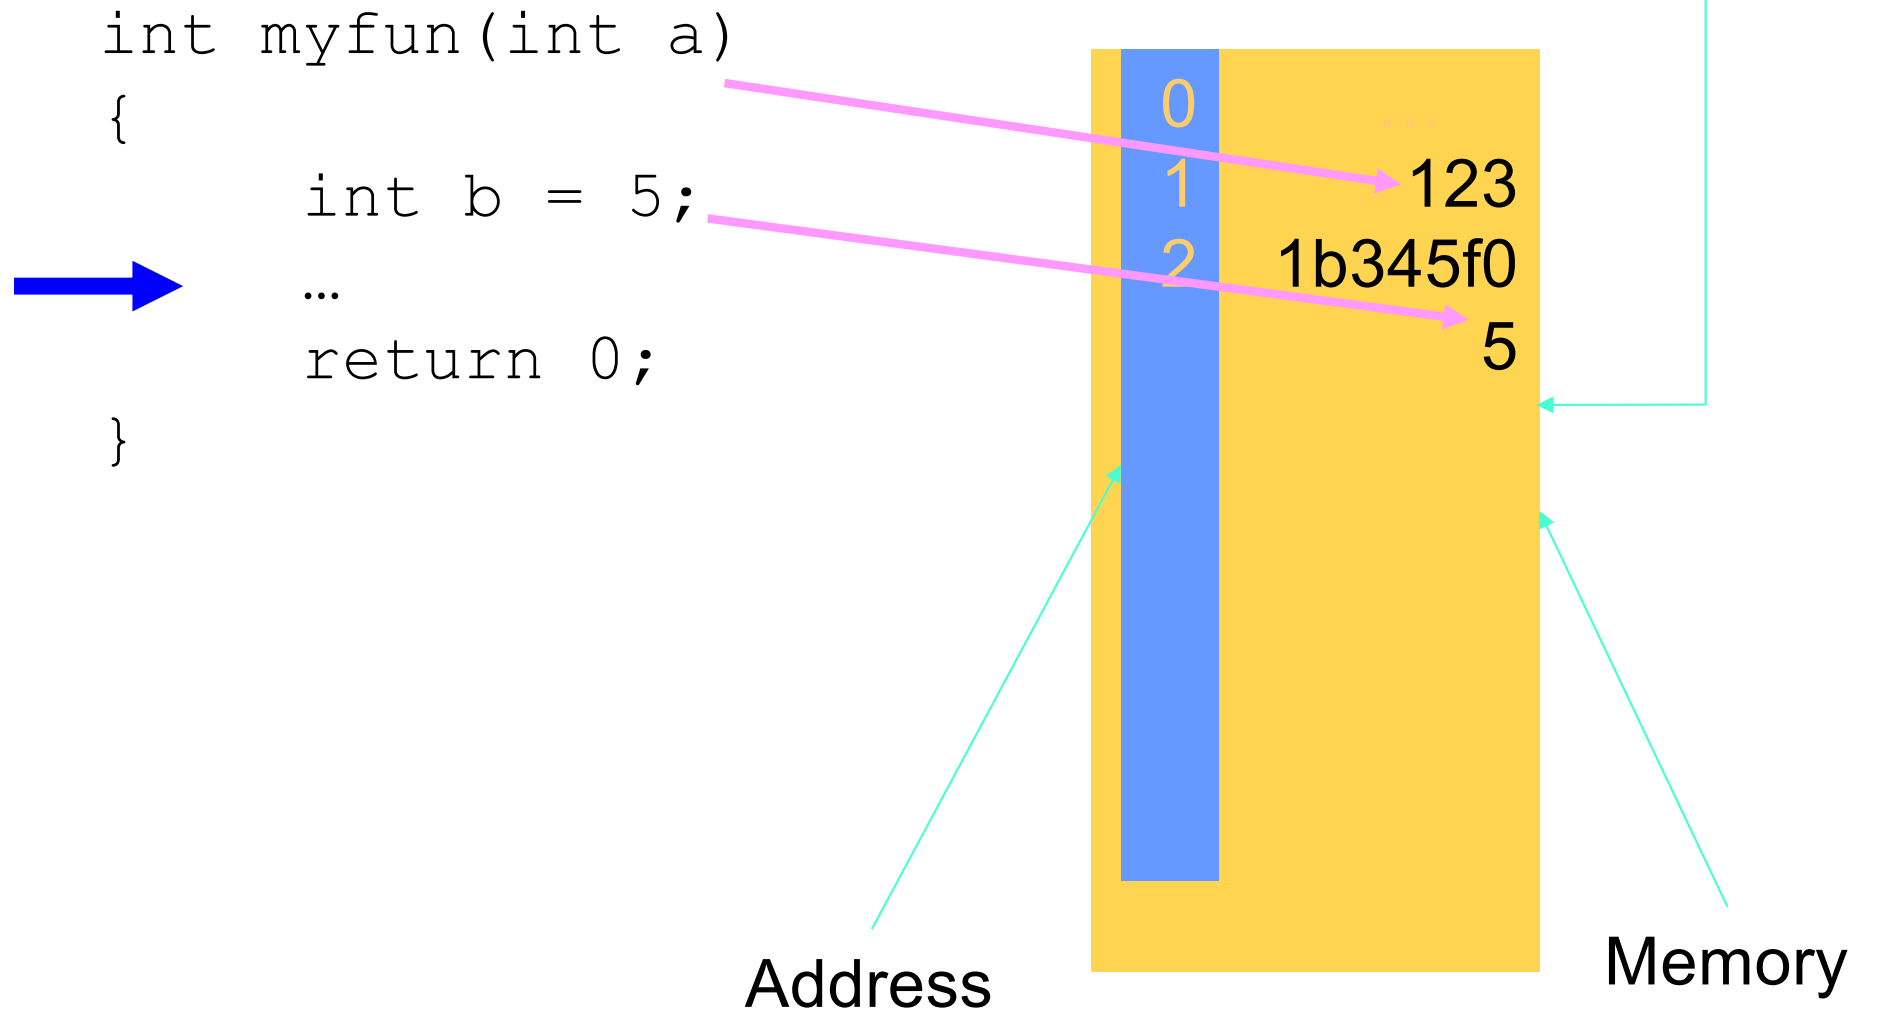


Function call example





Function call example





Function call example

```
int myfun(int a)
{
    int b = 5;
    ...
    return 0;
}
```

stack ptr

before retuning the value, all
the local variables are pop off

0
1
2

0

the return value is not fixed size

Address

Memory

Heap

Memory may be **dynamically** allocated at **run-time** from an area known as “the **heap**”.

Unlike the **stack**, which meets the temporary storage demands associated with called functions, the **heap** is accessed under direct programmer control.



We request an allocation of memory from the heap.

If there is sufficient **contiguous** memory available, we are given the address of the start of the allocated memory.

Pointer


will get back a pointer where first element of the byte of allocated memory is

used

free

newly
allocated

heap





Q: Where are parts of this program stored?

```
int a;    global
int main() {
    int b;    stack
    int *p;    stack    store an address
    p = malloc(...)    p points to heap
}
int doit(int c) {    stack
    static int d;    static
    means it persist foreve
}
```



```
int a;  
  
int main() {  
    int b;  
    int *p;  
    p = malloc(...)  
}  
  
int doit(int c) {  
    static int d;  
}
```

Q: What is the following *Java* code doing?

stack
`myObject fred = new myObject();`

heap - memory allocation

heap contains data
stack contains variables

A: Creating an object of type *myObject*.

However, what you *don't* see is the memory allocation required to instantiate the object.

Java also hides the act of freeing memory via automatic “garbage collection”.



SUMMARY

Memory allocation is ***not*** difficult!

It only causes problems because novice programmers may not recognise the **need** to address it...

Java programmers are less likely to experience such problems simply because *Java* hides the need to deal with this whole issue.



Memory Management Functions





Memory allocation functions

Memory allocation functions return a “**pointer to void**” . `void *`

A “**pointer to void**” is used to represent a pointer with no scalar value.

The pointer must therefore be cast to a specific type.



Memory allocation functions: malloc

```
#include <stdlib.h>
void *malloc(size_t size);
```

Typically defined as:
`typedef unsigned int size_t;`

Requests **size** number of bytes of memory.

Returns a pointer to the allocated memory, if successful, or a **NULL** pointer if unsuccessful

`size_t` is used for network, any kinds of binary read or write

A comment on the use of `size_t`:



Use of `size_t` replaces the use of more specific types, such as `int`, `short`, etc. This allows the actual implementation to be system-specific.

The `sizeof` operator is of type `size_t`. This is often used to specify memory requirements, so it makes sense to have the size argument in memory allocation functions of type `size_t`.



```
int * ptr; stack  
ptr = (int *)malloc(sizeof(int)*20);  
type casting how many bytes we want
```

If an **int** is 4 bytes, then this call will request 80 bytes of memory from the heap.



calloc

```
#include <stdlib.h>
void *calloc(size_t num, size_t size);
```

This is similar to `malloc` except that:

- It has two arguments:
 - `num` specifies the number of “blocks” of contiguous memory
 - `size` specifies the size of each block
- The allocated memory is cleared (set to ‘0’).



free

```
#include <stdlib.h>
void free(void *ptr);
```

ptr = malloc()
free(ptr)

```
#include <stdio.h>
#include <stdlib.h>

int main() {

    int *nums = NULL;

    nums = (int *) malloc(sizeof(int) * 10);
    int i;
    for (i = 0; i < 10; ++i) {
        nums[i] = 9990 + i;
        printf("%d: %d\n", i, nums[i]);
    }
    free(nums);
    nums = NULL;
    return 0;
}
```

This is used to de-allocate memory previously allocated by any of the memory allocation functions.



debug cmd
gcc file.c -o file -g
gdb ./file

cmd you can input:
r - reading
br - backtracing
print i
print nums[0]


```
int * ptr = NULL;  
ptr = (int *)malloc(sizeof(int)*20);
```

```
free((void *)ptr);
```

```
ptr = NULL;
```

good practice

ptr



realloc

1. allocate new size of memory
2. copy each of element into new area memory
3. if sucessful, remove the old area memory, assign to new area memory

```
#include <stdlib.h>
```

```
void *realloc(void *ptr, size_t size);
```

original pointer new size

This takes previously-allocated memory and attempts to resize it.

This may require a new block of memory to be found, so it **returns a new void pointer** to memory.

Contents are preserved.



ptr has changed

```
int * ptr;  
ptr = (int *)malloc(sizeof(int)*2);
```

```
ptr = (int *)  
    realloc(ptr, sizeof(int)*200);
```


ptr

```
char *text;  
text = (char *)malloc(sizeof(char) * 10);  
int i;  
for (i = 0; i < 6; ++i) {  
    text[i] = "catdog"[i];  
}  
text[i] = '\0';  
printf("text is %s\n", text); // text is catdog  
printf("address of text is %p\n", text);  
// address of text is 0x7f974e401870  
  
text = (char *) realloc(text, sizeof(char) * 10000);  
printf("text is %s\n", text); // text is catdog  
printf("address of text is %p\n", text);  
// address of text is 0x7f974e801600  
  
free(text);  
return 0;
```





```
int * ptr;  
ptr = (int *)malloc(sizeof(int)*2);
```

```
ptr = (int *)  
    realloc(ptr, sizeof(int)*200);
```



ptr



Dynamically creating structures

```
struct thing *    ptr;
```

```
ptr = (struct thing *)malloc(sizeof(struct thing));
```

```
/* Do stuff */
```

```
ptr->day = mon;
```

```
...
```

```
free((void *)ptr);
```

```
ptr = NULL;
```

This is a some of what *Java* does
“behind the scenes” on object
creation.



Safety issues



Safety issues

Caution #1:

De-allocate memory that is no longer required.

While the system should de-allocate resources on termination, it is **good practice** to take control of this process.

In some *Java* programs there is a noticeable performance dip when the automatic “garbage collection” functionality kicks in.

Safety issues

Caution #2:

wrapper function to malloc, keep track of if free or not

NEVER attempt to de-allocate memory that has not been allocated!

A common error is to try to free memory that has already been de-allocated, or was never allocated in the first instance.

Safety issues

Caution #3:

NEVER try to use memory that has been de-allocated.

This is also a common error leading to serious problems.



Safety issues

Caution #4: you may under estimate how many memory you need

Know your memory allocation requirements!

Use of the **sizeof** operator addresses the more obvious problems.

However, a common problem is to forget that a string includes a **'\0'** terminating character.

Safety issues

Caution #5:

Check for success!

A failed memory allocation request can lead to disaster if it is simply assumed to be successful.

Previous examples here have made this assumption for convenience. This would **NOT** qualify as bullet-proof code!

Safety issues

Typically, safe memory allocation is addressed by wrapping the relevant function in some additional code.

The following code^{*} demonstrates an example using **realloc**.

^{*} Adapted from Kay & Kummerfeld, *C Programming in a UNIX environment*

Safety issues

```
#include <stdlib.h>
```

```
void *
```

```
srealloc(void *ptr, size_t size)
```

Do error checking!

```
{
```

```
    void *res;
```

```
    if((res = realloc(ptr, size)) == (void *)0)
```

NULL

```
    {
```

```
        perror("realloc()");
```

print to console

```
        exit(1);
```

```
    }
```

```
    return res;
```

If the returned result is a NULL pointer, let the system print the appropriate error message via **perror** and then **exit**.

```
}
```

Otherwise, return the pointer to memory.



Safety issues

- **salloc**
- **srealloc**
- **Test return value on malloc, realloc**



Summary

- ✓ Understand the need for memory allocation and de-allocation
- ✓ Be able to use relevant C functions for achieving this
 - ✓ malloc
 - ✓ calloc not really
 - ✓ realloc
 - ✓ free
- ✓ Be able to allocate and access memory *sa*fely

Sources

- Image sources:
 - zazzle t-shirt
 - http://www.hazoment.com/Humor-Fasten_Safety_Belts.jpg
- Kay, J. & B. Kummerfeld (1989). *C Programming in the UNIX environment*. Addison-Wesley: Sydney.