
COMP2017 / COMP9017 Week 6 Tutorial

File IO, Function Pointers and Signals

File IO and Function Pointers

As per the unix philosophy, "*Everything is a file*", this means that we can typically get majority of our information from unix processes, pipes and memory mapped files.

Processes running inside UNIX contain a table of open files, known as the file descriptor table. The entries within the table are a simple numerical identifier that is known as a file descriptor. Each process normally contain three entries, 0, 1, 2, which represent standard input (stdin), standard output (stdout), and standard error (stderr). Any subsequently opened files are stored at the next available entry in the file descriptor table. When a process opens a file (or any other file-like object), the next free file descriptor is used. For example, if the file descriptor table contains the entries 0, 1, 2, 4, 6, the next opened file would take the value 3.

stdio.h functions and macros

By including the `stdio.h` header we have access to the `fopen`, `fwrite`, `fread`, etc. However, these functions are library functions for system the open, write, read, etc system calls. These calls utilise a `FILE*` object that contains platform specific information.

Using the `man` command, open up `stdio` man page like so: `man 3 stdio`

You can get a quick description of the function and how it is used.

fcntl.h andunistd.h functions and macros

`stdio` usually provides all we want for input and output, however we will need to occasionally use system calls to **handle the dirty work**. You may want to have **full control over an IO device** and specify the IO modes such as block and nonblocking, buffering modes, etc. To, open a file using a sys-call we can use the `open()` function which will return an integer (file descriptor) that corresponds to the handle attached to the process.

In relation to file descriptors, we have access to `read()` and `write()` system calls that operate on the file descriptors. These the the most primitive input/output functions that are commonly exposed on unix systems.

Write, Read and Flush

`fwrite` and `fread` are an abstraction that utilises binary stream input and output functions. Once a `FILE*` object has been initialised, you can perform file stream operations such as `fread` and `fwrite`.

These functions operate in a similar pattern where they require you **utilise a buffer for copying data and specifying the size of the data and number of elements**. This can be used in conjunction with any kind of pointer or array of any type assuming the bounds of the buffer supplied are not violated. You will need to acknowledge that you are reading and writing binary data and the structures you utilise will be interpreted based on your assumptions.

Usage example:

```
//Assume file has been opened for reading
...
int numbers[24];
size_t bytes_read = fread(numbers, sizeof(int), 24, f);
```

`fread` will return how many bytes have been read and will **move the file cursor** (where we are currently reading at) **forward 'sizeof(int)*24' bytes**.

```
int numbers[6] = {1, 2, 3, 4, 5, 6};
size_t bytes_written = fwrite(numbers, sizeof(int), 6, f);
```

Since it is file stream this may mean any contents you have written using `'fwrite'` **may not have been flushed or written to disk**. To immediately write to disk, you

```
//If you want data immediately flushed
fflush(f);
//if you have finished writing and want to close
fclose(f);
```

- For functions related to text formatting and printing, look into file stream counterparts for `gets`, `printf`, `scanf` and `getc`.

Blocking and Non-blocking IO

These two dictate how a function will behave when executing them in an IO context. **Blocking functions will wait until data in the buffer has been flushed to it before proceeding onto the next instruction.**

However, **non-blocking** will not wait for the data to be entered into the buffer and **will immediately check the buffer and grab the data it can**. This can seem counter-intuitive because someone may ask the question "Why would we want this if we can get incomplete data?". Typically non-blocking is used in **conjunction to signaling and interrupts**, for the signal to tell the program that the buffer is ready and for the program to read the buffer.

Flushing operates on an output stream (when we are writing), it will write all the data currently held in the buffer to its destination. Flushing an input stream is undefined behaviour.

Pre-Tutorial Questions

Question 1: Writing to standard out without printf

Knowing that processes have file descriptor table and the first 3 entries are for `stdin`, `stdout` and `stderr`, use the system calls to write to output "Hello World" using the `stdout` file descriptor.

Question 2: Changing standard output

There's nothing special about file descriptors 0, 1, or 2, besides the fact that these are where `stdin`, `stdout`, and `stderr` go. The `close()` system call can remove an entry from the file descriptor table, freeing up a number for reuse. The `dup()` (short for *duplicate*) system call makes a copy of a file descriptor into the lowest available index in the table. Using these together allows the programmer to replace what "standard out" is:

fd	Destination	fd	Destination	fd	Destination	fd	Destination
0	Terminal	0	Terminal	0	Terminal	0	Terminal
1	Terminal	1	Terminal	1	(empty)	1	fd
2	Terminal	2	Terminal	2	Terminal	2	Terminal
3	(empty)	3	(fd)	3	(fd)	3	fd
4	(empty)	4	(empty)	4	(empty)	4	(empty)
Program start		<code>open("filepath", O_WRONLY)</code>		<code>close(STDIN_FILENO)</code>		<code>dup(fd)</code>	

Table 1: The file descriptor table over the course of the program.

After the process table has been rearranged like this, any writes to standard out will actually go to the file instead. Our program will now be writing to the file, instead of the terminal, for its standard output.

Tutorial Questions

Question 3: File Descriptors

Discuss with your solution to the pre-tutorial work. Try and answer the following questions about file descriptors.

- What is a file descriptor and the file descriptor table
- What is a file and a stream
- What is non-blocking IO, how does it differ from blocking IO

1.

- file descriptor is a number that uniquely identifies an open file in a computer's operating system.
- file descriptors index into a per-process file descriptor table maintained by the kernel, that in turn indexes into a system-wide table of files opened by all processes

2.

- a stream is a generic interface for performing certain I/O operations.
- When you want to do input or output to a file, you have a choice of two basic mechanisms for representing the connection between your program and the file: file descriptors and streams

3.

- Well blocking IO means that a given thread cannot do anything more until the IO is fully received (in the case of sockets this wait could be a long time).
- Non-blocking IO means an IO request is queued straight away and the function returns. The actual IO is then processed at some later point by the kernel.
- Blocking - Linear programming, easier to code, less control.
- Non-blocking - Parallel programming, more difficult to code, more control.

- What is a file descriptor?
a handle used to access a file or other I/O resource. int type
- What is the file descriptor table?
a table maintained by the OS to keep track of file descriptors
- What is the difference between a file and a stream?
File - int type
Stream - FILE* type

ssize_t

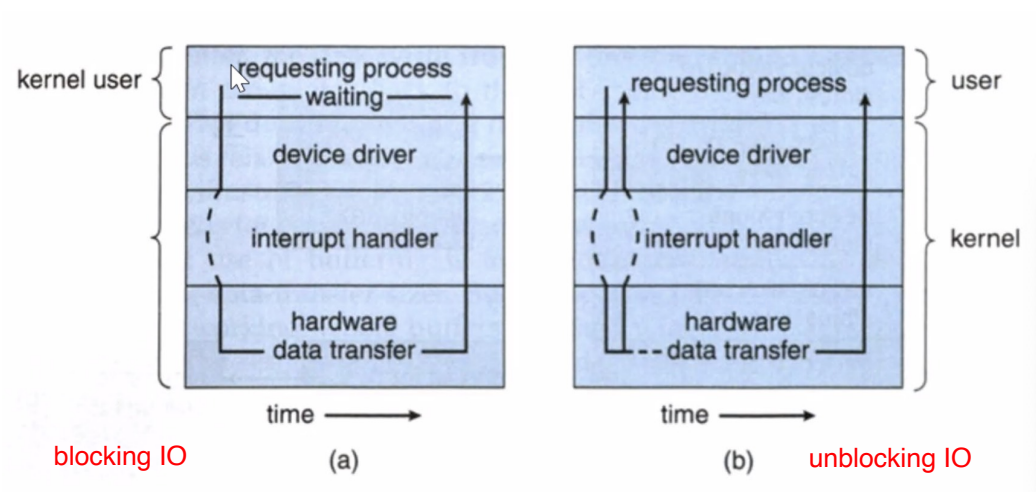
read(int fildes, void *buf, size_t nbyte)

- system call
- unbuffering read

size_t

fread(void *restrict ptr, size_t size, size_t nitems, FILE *restrict stream)

- c abstraction for file
- buffered read



Function pointers

Function pointers are another type within C that allows for variables to hold an address to a function. These allow you to have a pointer variable refer to a function and execute it. The function signature forms part of the type information for the variable which allows the compiler to confirm that an appropriate function with the same signature has been assigned.

Example:

```
int add(int a, int b) {
    return a + b;
}

int main() {
    int (*a)(int, int) = &add;
    //We have set the function pointer a to the address of add
    int result = a(1, 3);
    printf("Result of add: %d\n", result);
}
```

It is common to utilise function pointers as a parameter to a function for generalising certain operations on a collection and to act as a callback. Function pointers can exist within struct fields, although this is considered to be bad practice and an attempt to employ a simplistic object oriented pattern in C.

Example:

```
int calc(int a, int b, int (*c)(int, int)) {
    return c(a, b);
}

int add(int a, int b) {
    return a + b;
}

int main() {
    //We passed add to the calc function,
    //if we have other functions that have a similar signature

    int result = calc(1, 3, &add);

    printf("Result of add: %d\n", result);
}
```

- What would the pros and cons be of using typedef on a function pointer?
- Why would we pass a function pointer to another function? function pointer are also useful for callback functions
- Has there been a programming language where you have used something similar?

function reference in java

Question 4: Sorting and outputting

Extend the following code to implement bubble sort. Bubble sort is an inefficient but simple to implement sorting algorithm. It compares each pair of adjacent elements on each iteration, eventually bubbling largest (or smallest) to their ordered position.

The algorithm swaps the pair if they are in the wrong order (This is determined by the `cmp` function). This is repeated until the array is sorted.

```
int element_cmp(void* a, void* b);

void bubble_sort(void* elements, size_t n_elements,
                 size_t size_element, int (*cmp)(void*, void*));
```

Output the contents of the sorted list afterwards to verify the elements have been sorted. Try this sorting algorithm on the following types:

- Integers, `int`
- Strings, `char*`
- `struct scorecard char name[20]; int score;`

The following is the pseudo code for bubble sort.

```
fn bubblesort(list, length):
  let i = 0
  while i < length:
    let j = 0
    while j < length - 1:
      if compare(list[ij], list[j+1] > 0):
        swap(list[j], list[j + 1])
      j = j + 1
    i = i + 1
end def
```

Question 5: Replacing Bubblesort

Obviously there are better sorting algorithms than bubble sort. The C standard library (`stdlib.h`) includes sorting algorithm of its own that you can use. You should be able to identify the similarities between the bubble sort function and the `qsort` function and easily replace your bubble sort with `qsort`.

Use the man pages to retrieve `qsort`'s function signature.

Question 6: Tiny delegator

You are given a list of string instructions to parse. You are required to build a list of these instructions that map to a basic function pointer and value. This will represent a very simple program that calculates

After the instructions have been created and added to the list, your program will execute list of functions.

Specification:

ADD - Addition
SUB - Subtraction
MUL - Multiple
DIV - Divide

Your program should implement these 4 operations as well as parse each line, your program only has to deal with fixed 32 bit integer numbers (not floating point).

The instructions will be provided in this form:

```
ADD 9 10
SUB % 10
MUL 3 %
DIV % 1
```

The % symbol represents the value from the previous operation, and will need to be referred to by that location.

- a form of software interrupt
- allows one process to communicate with another

Signals

When you press `ctrl+c` or `ctrl+z`, your shell sends the signal **SIGINT** or **SIGSTOP** respectively to the current fore-ground program. Signals are a way to communicate with programs that are currently running. You can send a signal to a program if you know its process ID using the misleadingly named `kill` command.

Give a process we want to send a signal to has the process id of 255, we can send a signal using the `kill` command.

Format

```
kill -<signal type> <pid>
#example
kill -SIGKILL 255
```

type `ps` to get pid

Other useful signals to know are **SIGTERM**, which **asks** the process to terminate, and **SIGKILL** which **forces** the process to terminate. All signals apart from **SIGKILL** and **SIGSTOP** can be handled by the process by defining a signal handler.

Custom signal handling

We can setup custom signal handlers to our programs that allow us to setup functions to trigger for certain signals. On the event of a signal being received by the process, it will map the signal to a function, this will interrupt the current execution and start executing the current function binding for that signal.

Example:

```
void interrupt_handler(int sig) {
    printf("I was interrupted!\n");
}

int main() {
    signal(SIGINT, interrupt_handler);
    ... rest of code
}
```


However, the `signal` function has been deprecated in favor of `sigaction` which setup is a little less obvious.

```
sigaction(int sig, const struct sigaction *restrict act, struct sigaction *restrict oact);  
  
void sigint_handler(int signo, siginfo_t* sinfo, void* context) {  
    printf("I was interrupted\n");  
}  
  
int main() {  
    struct sigaction sig;  
    memset(&sig, 0, sizeof(struct sigaction));  
    sig.sa_sigaction = sigint_handler; //SETS Handler  
    sig.sa_flags = SA_SIGINFO; void* memset(void *b, int c, size_t len);  
                                writes len bytes of value c to the string b.  
  
    if(sigaction(SIGINT, &sig, NULL) == -1) {  
        perror("sigaction failed");  
        return 1;  
    }  
}
```

You can get information about `sigaction` function and types from the man pages.

Question 7: A custom signal handler

The C standard library provides the `signal.h` header for handling signals.

Compile and run the following program:

```
#include <stdio.h>  
#include <signal.h>  
#include <unistd.h>  
#include <stdbool.h>  
  
struct sigaction {  
    void (*sa_handler)(int);  
    void (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_restorer)(void);  
};  
  
int main(void) {  
  
    signal(SIGINT, SIG_IGN);  
  
    while (true) {  
        sleep(1);  
    }  
  
    return 0;  
}
```

SIGINT - signal interrupt
SIG_IGN - ignore signal

1. control + Z
2. kill -<signal type> <pid>

Attempt to kill the program via `SIGINT`. Find a way to kill the program.

Hint: You may want to use the `kill` command from another terminal. Alternatively, you could kill the program within the same terminal.

Question 8: Tell me the time!

After discovering that you can assign a function to a signal handler. Write a program that will wait for signal (specifically `SIGUSR`) to be sent to the process and output the current time the signal was received. Your program must *gracefully* exit when it receives `SIGINT`.

```
$ ./handle
Thu Apr 15 12:22:34 2018          SIGUSR - user defined signal
Thu Apr 15 12:23:12 2018
Thu Apr 15 12:25:41 2018
^C
Shutting down
```

Hint: You will need to use the `time.h` header and the utilise the `time_t` and `time` function.

Question 9: Signal Assistant

You are to construct a program that will simply respond to signals. Instead of commanding it via standard input this program will respond to your signals.

- When `SIGINT` is sent, the process should output the number minutes the computer has been on, check `/proc/uptime` use “`sysctl kern.boottime`” to get uptime in macOS
- When `SIGUSR1` is sent, the process should provide a random number from `/dev/urandom`
- When `SIGUSR2` is sent, the process should tell the user a joke
- When `SIGHUP` is sent, tell the user that they are going home and cannot respond anymore